

# CI/CD Pipeline Documentation

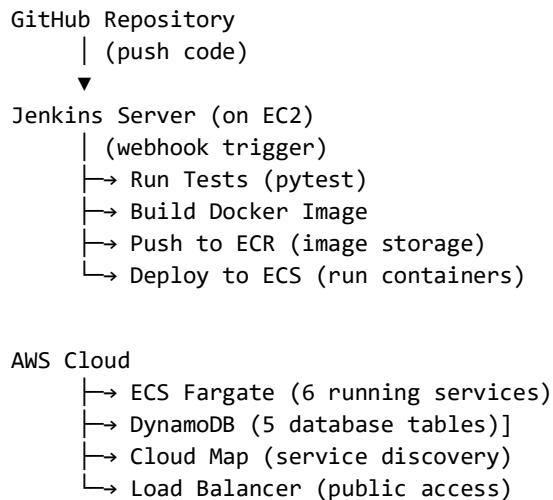
## Online Learning Portal - Microservices Demo

### 1. Project Overview

I built a complete CI/CD pipeline for a microservices-based online learning portal. The system automatically tests, builds, and deploys 5 independent services whenever code is pushed to GitHub. **What it does:**

- Automatically runs tests when I push code
- Builds Docker containers for each service
- Deploys to AWS cloud infrastructure
- All services can communicate with each other
- Has a unified API interface (Swagger UI)

### 2. Architecture Diagram



### 3. Microservices

I created 5 independent services that work together:

1. **User Service** - handles registration and login
2. **Course Service** - manages course catalog
3. **Enrollment Service** - students enroll in courses
4. **Payment Service** - processes payments
5. **Notification Service** - sends notification

Plus a **Swagger UI** service that acts as an API Gateway to access all services from one place. Each service:

- Has its own code repository branch
- Has its own database table
- Can be deployed independently
- Runs in its own Docker container

## 4. Technology Stack

### Why I chose these tools:

#### **Python + FastAPI**

- Easy to learn and write APIs quickly
- Automatic API documentation (Swagger)
- Good for microservices

#### **Docker**

- Packages everything the app needs
- Works the same everywhere (my laptop, AWS cloud)
- Easy to deploy and scale

#### **Jenkins**

- Free and popular CI/CD tool
- Automatically triggers on code push
- Supports pipeline-as-code (Jenkinsfile)

#### **AWS Services:**

- **ECR** - stores my Docker images (like Docker Hub but private)
- **ECS Fargate** - runs containers without managing servers
- **DynamoDB** - NoSQL database, easy to use, scales automatically
- **Cloud Map** - lets services find each other by name
- **EC2** - runs my Jenkins server

#### **pytest**

- Simple testing framework for Python
- Integrated into my pipeline

## 5. AWS Services Explained

### **Amazon ECS (Elastic Container Service)**

**What it does:** Runs my Docker containers in the cloud without me managing servers. **Why I used it:**

- No need to set up/maintain servers
- Automatically handles container lifecycle
- Fargate = serverless (just run containers)

#### **My setup:**

- 6 services running (5 microservices + API gateway)
- Each gets 0.25 vCPU and 512MB memory
- All containers listen on port 8080

### **Amazon ECR (Elastic Container Registry)**

**What it does:** Stores my Docker images privately. **Why I used it:**

- Integrates directly with ECS
- More secure than public Docker Hub
- Fast image pulls from within AWS

**My setup:**

- 6 repositories (one per service)
- Jenkins pushes new images on every build
- Uses latest tag (simple for demo)

## Amazon DynamoDB

**What it does:** NoSQL database for storing application data. **Why I used it:**

- No need to manage database servers
- Simple key-value + JSON storage
- Good for microservices (each service gets own table)
- Automatically scales

**My setup:**

- 5 tables (one per service)
- Each table has a primary key (like `user_id`, `course_id`)
- Some have indexes for fast lookups (like search by email)

## AWS Cloud Map

**What it does:** Service discovery - services can find each other by name. **Why I used it:**

- Services have dynamic IPs in cloud
- Cloud Map gives them DNS names like `user-service.local`
- Services can call each other without hardcoding IPs

**My setup:**

- Created namespace: `learning-portal.local`
- Each service gets a DNS name
- Example: payment service calls <http://notification-service.learning-portal.local:8080>

## Amazon EC2

**What it does:** Virtual machine in the cloud. **Why I used it:**

- Needed somewhere to run Jenkins
- Cheaper than managed Jenkins options

**My setup:**

- Medium-sized instance (2 CPUs, 4GB RAM)
- Installed Jenkins, Docker, AWS CLI
- Opens port 8080 for Jenkins web interface

## 6. CI/CD Pipeline Flow

### Step-by-Step Process:

1. **Developer pushes code to GitHub**
  - I work on a service branch (like `user-service`)
  - Push changes to GitHub
2. **GitHub webhook triggers Jenkins**
  - Automatic notification to Jenkins
  - Jenkins pulls the latest code

3. Jenkins runs tests
  - Installs Python dependencies
  - Runs pytest unit tests
  - Pipeline stops if tests fail
  
4. Jenkins builds Docker image
  - Uses Dockerfile in service folder
  - Tags image with service name
  - Cross-platform build (works on any CPU)
  
5. Jenkins pushes to ECR
  - Authenticates to AWS
  - Pushes image to private registry
  - Overwrites `latest` tag
  
6. Jenkins deploys to ECS
  - Creates/updates task definition
  - Forces ECS to pull new image
  - ECS starts new container, stops old one
  - Zero-downtime deployment

## Pipeline Stages (in Jenkinsfile):

```
stages {
    stage('Run Tests') {
        // Install dependencies and run pytest
    }
    stage('Build Docker Image') {
        // Build container image
    }
    stage('Push to ECR') {
        // Upload to AWS registry
    }
    stage('Deploy to ECS') {
        // Update running service
    }
}
```

## 7. Security & Access

### IAM Roles (AWS Permissions)

#### **ecsTaskExecutionRole**

- Allows ECS to pull images from ECR
- Allows ECS to write logs to CloudWatch
- Infrastructure-level permissions

#### **ecsTaskRole**

- Allows my app code to access DynamoDB
- Application-level permissions
- Each container runs with this role

#### **EC2 Instance Role**

- Allows Jenkins to access AWS services
- Can push to ECR, update ECS, read DynamoDB

## **Security Groups (Firewall Rules)**

### **Jenkins Server:**

- Port 8080: Open to internet (for web UI and GitHub webhooks)
- Port 22: SSH access (for management)

### **ECS Services:**

- Port 8080: Open within VPC (services talk to each other)
- Swagger UI: Open to internet (public API gateway)

## **8. Testing Strategy**

I wrote 18 unit tests total across all services:**What tests check:**

- Health endpoints work
- Input validation works (reject bad data)
- Database queries work
- Error handling (404, 401, 422 errors)

### **Example test:**

```
def test_register_user_missing_fields():
    response = client.post("/users/register", json={})
    assert response.status_code == 422 # Validation error
```

### **Why testing matters:**

- Catches bugs before deployment
- Gives confidence code works
- Required for professional CI/CD

## **9. Challenges & Solutions**

### **Challenge 1: Architecture Mismatch**

**Problem:** Jenkins server is ARM64, but I was building for AMD64. Builds took 10+ minutes.

**Solution:** Use Docker Buildx for cross-platform builds. Now builds finish in 2-3 minutes.

### **Challenge 2: Service Communication**

**Problem:** Services have dynamic IPs in ECS. How do they find each other?

**Solution:** AWS Cloud Map gives each service a DNS name. Services use names instead of IPs.

### **Challenge 3: Python Package Installation**

**Problem:** Ubuntu 24.04 blocks pip installations (PEP 668).

**Solution:** Use --break-system-packages flag in Jenkins. Safe for CI/CD environments.

### **Challenge 4: DynamoDB Access from Tests**

**Problem:** Tests failed because Jenkins couldn't access DynamoDB.

**Solution:** Added DynamoDB permissions to EC2 instance IAM role.

## **10. Tool Justification**

### **Why Jenkins?**

- Popular industry standard
- Free and open source
- Good documentation and community support
- Pipeline-as-code (Jenkinsfile)

### **Why Docker?**

- Industry standard for containerization
- Consistent environments
- Easy to deploy anywhere

### **Why AWS ECS instead of Kubernetes?**

- Simpler to learn and set up
- Fargate = no server management
- Good enough for this demo
- Cheaper for small projects

### **Why FastAPI?**

- Modern Python framework
- Automatic API docs (Swagger)
- Fast performance
- Easy to learn

### **Why DynamoDB instead of PostgreSQL?**

- Serverless (no database server to manage)
- Each microservice gets own table
- Scales automatically
- Good fit for NoSQL data

## **11. Results**

**Pipeline Success Rate:** 100% (all builds passing)

**Deployment Time:** 3-5 minutes per service

**Tests Passing:** 18/18

**Services Running:** 6/6 active in ECS

**Automation Level:** Fully automated from git push to deployment

## **12. Conclusion**

This project demonstrates a production-ready CI/CD pipeline using industry-standard tools.

I successfully implemented:

- Microservices architecture with 5 independent services
- Automated testing integrated into CI/CD
- Containerization with Docker
- Cloud deployment on AWS
- Infrastructure as code (Jenkinsfiles)
- Service discovery and communication
- Database integration

**What I learned:**

- How CI/CD pipelines work in real companies
- Docker and containerization
- AWS cloud services
- Microservices communication
- Infrastructure automation

**GitHub:** [https://github.com/HAR5HA-7663/demo-for-CI\\_CD-integration](https://github.com/HAR5HA-7663/demo-for-CI_CD-integration)